

Demo: On working with dedicated SQL pool, and how we can create data, insert data or analyze data in SQL pool.

First I will create a table in SQL pool (Taxipool), next I will create folder called NYC Taxi, you can create a script in this folder or you can import a script from your desktop as well. I have imported it from my desktop.

This table is a Round robin distribution, basically Round robin means to evenly distribute data into different distributions of table.

```
1 /*
2 This script will finish in around 60 seconds. It loads 2 million rows of NYC Taxi data into a table called dbo.Trip
3 */
4 CREATE TABLE [dbo].[Trip]
5 (
6     [DateID] int NOT NULL,
7     [MedallionID] int NOT NULL,
8     [HackneyLicenseID] int NOT NULL,
9     [PickupTimeID] int NOT NULL,
10    [DropoffTimeID] int NOT NULL,
11    [PickupGeographyID] int NOT NULL,
12    [DropoffGeographyID] int NOT NULL,
13    [PickupLatitude] float NULL,
14    [PickupLongitude] float NULL,
15    [PickupLatLong] varchar(50) COLLATE SQL_Latin1_General_CP1_CI_AS NULL,
16    [DropoffLatitude] float NULL,
17    [DropoffLongitude] float NULL,
18    [DropoffLatLong] varchar(50) COLLATE SQL_Latin1_General_CP1_CI_AS NULL,
19    [PassengerCount] int NULL,
20    [TripDurationSeconds] int NULL,
21    [TripDistanceMiles] float NULL,
22    [PaymentType] varchar(50) COLLATE SQL_Latin1_General_CP1_CI_AS NULL,
23    [FareAmount] money NULL,
24    [SurchargeAmount] money NULL,
25    [TaxAmount] money NULL,
26    [TipAmount] money NULL,
27    [TollsAmount] money NULL,
28    [TotalAmount] money NULL
29 )
30 WITH
31 (
32     DISTRIBUTION = ROUND_ROBIN,
33     CLUSTERED COLUMNSTORE INDEX
34 );
35
```

So basically here what we are doing, we are copying data into a Trip table from the open dataset. This script will insert a few million records into Trip table.

```
27 [TollsAmount] money NULL,
28 [TotalAmount] money NULL
29 )
30 WITH
31 (
32     DISTRIBUTION = ROUND_ROBIN,
33     CLUSTERED COLUMNSTORE INDEX
34 );
35
36 COPY INTO [dbo].[Trip]
37 FROM 'https://nytaxiblob.blob.core.windows.net/2013/Trip2013/QID6392_20171107_05910_0.txt.gz'
38 WITH
39 (
40     FILE_TYPE = 'CSV',
41     FIELD_TERMINATOR = '|',
42     FIELDQUOTE = '',
43     ROW_TERMINATOR = '0x0A',
44     COMPRESSION = 'GZIP'
45 )
46 OPTION (LABEL = 'COPY : Load [dbo].[Trip] - Taxi dataset');
```

Now I should be able to see some data into Trip table here.

The screenshot shows a data tool interface with a sidebar on the left containing 'Data', 'Develop', 'Integrate', 'Monitor', and 'Manage'. The main workspace is divided into a 'Workspace' pane on the left and a 'Query editor' on the right. The 'Workspace' pane shows a tree view of 'Databases' with 'Taxipool (SQL)' selected, containing 'Tables', 'External tables', 'Views', 'Programmability', 'Schemas', and 'Security'. The 'Query editor' shows a SQL query:

```
1 SELECT TOP (100) [DateID]
2 , [MedallionID]
3 , [HackneyLicenseID]
4 , [PickupTimeID]
5 , [DropoffTimeID]
6 , [PickupGeographyID]
7 , [DropoffGeographyID]
8 , [PickupLatitude]
9 , [PickupLongitude]
10
```

 The 'Results' pane shows a table with 10 columns: DateID, MedallionID, HackneyLicenseID, PickupTimeID, DropoffTimeID, PickupGeographyID, DropoffGeographyID, PickupLatitude, PickupLongitude, and PickupLatLong. The table contains 10 rows of data. The status bar at the bottom indicates '00:00:02 Query executed successfully.'

DateID	MedallionID	HackneyLicenseID	PickupTimeID	DropoffTimeID	PickupGeographyID	DropoffGeographyID	PickupLatitude	PickupLongitude	PickupLatLong
20130809	1872	39761	76619	76819	250384	12361	40.8012	-73.9454	40.8012; -73.9454
20131012	3835	1141	81540	82020	22168	127571	40.7298	-73.9836	40.7298; -73.9836
20131220	1193	38952	41493	42464	238602	33307	40.7297	-73.9932	40.7297; -73.9932
20130702	1632	25835	81780	82140	279919	91515	40.7445	-74.0065	40.7445; -74.0065
20130921	9156	4632	62562	63080	262615	92380	40.7343	-74.0061	40.7343; -74.0061
20130717	1027	30908	62300	62880	261643	57596	40.7799	-73.956	40.7799; -73.956
20130809	13190	41345	72265	72893	270648	30714	40.751	-73.9715	40.751; -73.9715
20131114	2842	18453	3989	4540	279919	250711	40.7443	-74.0067	40.7443; -74.0067
20130311	1551	40337	74412	74620	73134	301330	40.7258	-74.004	40.7258; -74.004
20131019	4441	36790	75900	76380	193394	40851	40.8028	-73.9637	40.8028; -73.9637

I have almost 2.8 million records in Trip table now.

The screenshot shows a data tool interface with a sidebar on the left containing 'Home', 'Data', 'Develop', 'Integrate', 'Monitor', and 'Manage'. The main workspace is divided into a 'Workspace' pane on the left and a 'Query editor' on the right. The 'Workspace' pane shows a tree view of 'SQL scripts' with 'SQL script 1' selected. The 'Query editor' shows a SQL query:

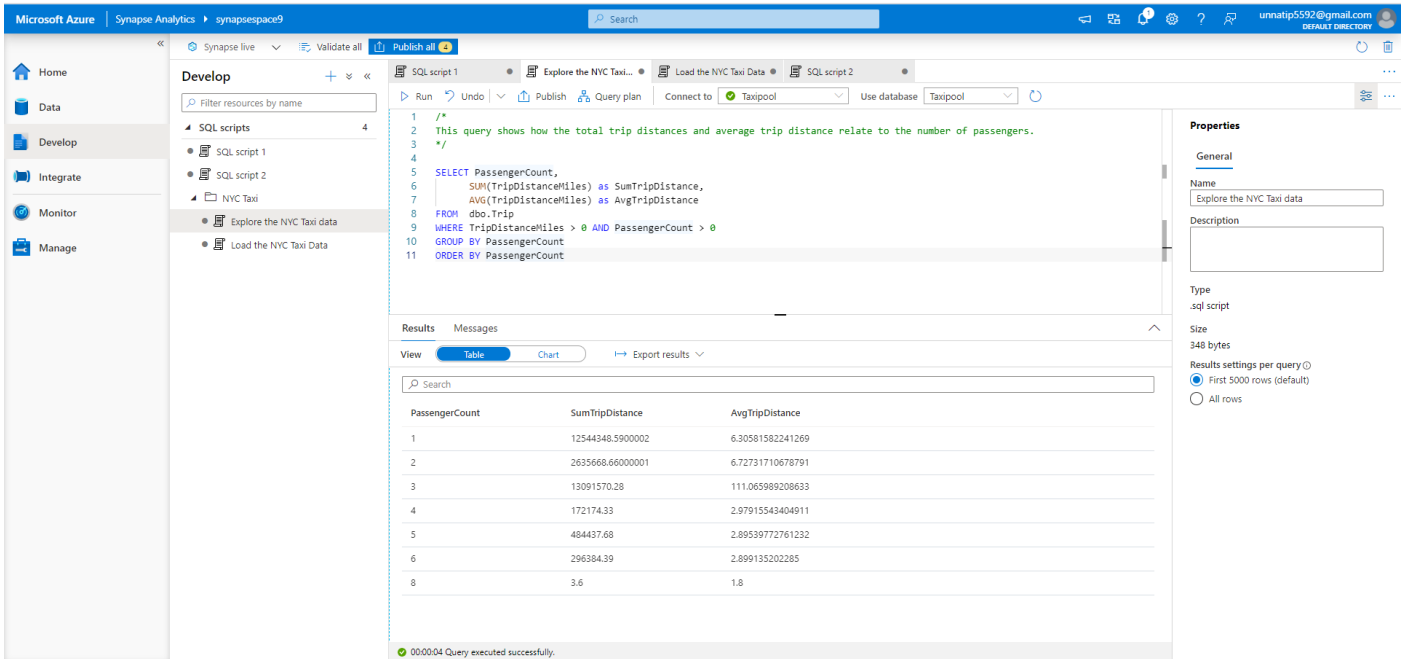
```
1 select COUNT(1) from dbo:trip
```

 The 'Results' pane shows a table with 1 column: COUNT(1). The table contains 1 row of data: 2838927. The status bar at the bottom indicates '00:00:01 Query executed successfully.'

COUNT(1)
2838927

let's explore some data from this table now in this query using a trip table, I'm taking some relevant data, data where the distance in miles is greater than Zero and and passenger count is greater than you.

This query will show the total trip distance and average trip distance related to number of passenger.



You can also convert this query result in a chart as well.

